

# OpenACC Tutorial



# What is OpenACC

- **Open Accelerators**
- Through various **compiler directives** to write GPU code
- Lower the technical barriers to GPU programming

# What is OpenACC

```
#pragma acc data copy(A) create(Anew)
while ( error > tol && iter < iter_max ) {
    error = 0.0;
    #pragma acc kernels
    {
        #pragma acc loop independent collapse(2) reduction(max:error)
        for ( int j = 1; j < n-1; j++ ) {
            for ( int i = 1; i < m-1; i++ ) {
                Anew [j] [i] = 0.25 * ( A [j] [i+1] + A [j] [i-1] +
                                         A [j-1] [i] + A [j+1] [i]);
                error = max ( error, fabs (Anew [j] [i] - A [j] [i]));
            }
        }
        ...
    }
}
```

<https://www.openacc.org/>

# CUDA

- CudaMalloc(...): Declare memory on the GPU
- CudaMemcpy(...): Move data
- functionName<<<thread, blocks>>>(...): Write your own Cuda

Kernel Function

=> High entry barrier

# OpenACC

- No need to declare the memory on the device
- `#pragma acc data copy(...)`: You can move data with a simple clause
- You can directly use parallel region to port to the GPU.

=> Easy to use

# OpenACC Directive

- #pragma acc <directive> <clause>
  - #pragma is a compiler hint
  - acc tells the compiler that this is the OpenACC pragma
  - **directive** is what OpenACC tells the compiler to indicate
  - **clause** is an instruction for OpenACC to supplement or optimize the directive.


# OpenACC Directive

- #pragma acc **parallel**
- parallel tells the compiler that this code should be **redundantly parallelized**

```
#pragma acc parallel  
{
```

When encountering the **parallel** directive, the compiler will generate 1 or more **parallel gangs**, which execute redundantly.

```
}
```




```
#pragma acc parallel  
{
```

```
for(int i = 0; i < N; i++)  
{  
    // Do Something  
}
```

```
}
```

This loop will be **redundantly parallelized** across the **gangs**

This means that each **gang** will execute the entire loop



# OpenACC Directive

- #pragma acc **parallel** **loop**
- **loop** tells the compiler that this loop needs to be parallelized
- It also tells the compiler that this loop can be safely parallelized.


```
#pragma acc parallel
{

    #pragma acc loop
    for(int i = 0; i < N; i++)
    {
        // Do Something
    }

}
```

The **loop** directive informs the compiler which loops to parallelize.

The iterations of the loop will be broken up evenly among the parallel **gangs**.



The **gangs** will then execute in parallel with one another.



# OpenACC Directive

- #pragma acc **parallel** **loop** **reduction**(<operation>:<target>)
  - **reduction** tells the compiler that a target is to be reduced
  - reduce: perform global operations on the selected target

```
int sum = 0;
```

```
#pragma acc parallel loop reduction(+:sum)
```


```
for(int i = 0, i < N, i++) sum += i;
```

# OpenACC Directive

- #pragma acc **kernels**
  - All actions are decided by the compiler
  - You can also include the sequential code

```
#pragma acc kernels
{
    for(int i = 0; i < N; i++)
    {
        // Do Something
    }
    for(int i = 0; i < M; i++)
    {
        // Do Something Else
    }
}
```

With the **kernels** directive, the **loop** directive is implied.



# OpenACC Directive

- #pragma acc **kernels** loop independent
  - Tell the compiler that this loop can be safely parallelized, and force parallelization of it

```
#pragma acc kernels  
{
```


```
  for(int i = 0; i < N; i++)  
  {  
    // Do Something  
  }
```

```
  for(int i = 0; i < M; i++)  
  {  
    // Do Something Else  
  }
```

```
}
```

This process can happen  
multiple times within the  
**kernels** region.

Each loop can have a different number of  
gangs, and those gangs can be  
organized/optimized completely differently.



# Data Management

You can copy only part of the data  
`#pragma acc parallel loop copy(A[1:N-2])`

- `#pragma acc data copy(...)`
  - Copy the data into the GPU and copy the data back to the CPU after the parallel region ends
- `#pragma acc data copyin(...)`
  - Copy the data into the GPU and delete the data on the GPU after the parallel region ends.
- `#pragma acc data copyout(...)`
  - Copy the data back to the CPU and delete the data on the GPU after the parallel region ends.
- `#pragma acc data create(...)`
  - Declare a space on the GPU without performing any copying operations
  - When there are variables for temporary storage, using this clause eliminates the need to copy in and out.

# Data Management

```
#pragma acc data copy(A[0:N])
```

```
#pragma acc parallel
```

```
{
```

```
    #pragma acc loop
```

```
    for(int i = 0; i < N; i++) A[i] = 0;
```

```
}
```



```
#pragma acc kernels copy(a[0:N])  
for(int i = 0; i < N; i++){  
    a[i] = 0;  
}
```

# Loop Optimization

- #pragma acc parallel loop **collapse(...)**
  - Can be used in tightly nested loops
  - collapse can flatten loops and turn multiple loops into one large parallel loop

```
#pragma acc parallel loop collapse( 2 )
```

```
for(int j = 0; j < M; j++) {  
    for(int k = 0; k < Q; k++) {  
        < loop code >  
    }  
}
```

TIP1:

When the outer loop is too small, flattening the loop can increase GPU usage.

# Loop Optimization

- #pragma acc parallel loop tile(x, y)
  - Calculate loop break for multiple tiles (blocks)

```
#pragma acc parallel loop tile( 32, 32 )
```

```
for(int j = 0; j < 128; j++) {  
    for(int k = 0; k < 128; k++) {  
        < loop code >  
    }  
}
```

TIP1:

Try to make the tile size a multiple of 32. The threads in a worker and vector of Nvidia GPU are executed in units of 32.

TIP2:

Do not use tiles larger than 32\*32, because in NVIDIA GPU, the maximum number of threads in a gang is 1024 (32\*32)